Note: No Week 11 slides were in the project — reconstructed from the assignment Q&A + the OECD AI Principles, EU AI Act, and the SaGE consistency paper. Q1–Q11 captured from the assignment; Q12 wasn't in the screenshots (flagged at the end).

1 Data-Subject Rights (recap)

Right to be forgotten (Q1): the data-privacy right that **requires companies to delete user data when requested** (GDPR Article 17). Distinct from right to data portability (get a copy), right to transparency (know how data is used), and IP retention.

2 OECD AI Principles

The OECD defines value-based principles for trustworthy AI. Key ones tested:

Principle	What it covers
Inclusive growth, sustainable development & well-being	AI should benefit people & the planet
Human-centred values & fairness	Respect rule of law, human rights, democratic values
Transparency & explainability	People can understand & challenge AI outcomes
Robustness, security & safety (Q2)	AI must **operate safely & function properly under normal use, misuse, or adverse conditions** throughout its lifecycle
Accountability (Q7)	AI actors are responsible; ensure **traceability** — trace back how data was collected & how it influences decisions

Assignment: 'operate safely under normal use, misuse or adverse conditions' = **Robustness, security & safety** (Q2). 'Traceability' under Accountability **ensures one can trace back how data was collected & how it influences decision-making** (Q7).

3 The EU AI Act — risk tiers

The EU AI Act is **horizontal, risk-based** regulation with four tiers: **Unacceptable** (banned), **High-risk** (strict obligations), **Limited-risk** (transparency duties), **Minimal-risk** (free).

Unacceptable-risk — banned (Q3, select all):

- **Social Scoring Systems**
- **Cognitive Behavioural Manipulation**

(Critical Infrastructure Management & Law Enforcement are **high-risk**, not banned.)

Deepfake / generated-media transparency (Q5): the EU AI Act requires AI-generated media to be **clearly labelled as AI-generated** (a limited-risk transparency duty).

The 'deployer' in the AI value chain (Q6): the **person or organization that uses the AI system under its own authority** (distinct from provider/developer, importer, or regulator).

4 Threat-Mitigation Styles

Silently eliminating the threat (Q4): when a system (like an email spam filter) removes malicious content **in the background without user intervention** — as opposed to educating users, showing interface-level warnings, or enforcing legal rules. It just quietly handles it.

5 SaGE — Semantic Graph Entropy (LLM moral consistency)

SaGE is a framework to measure an LLM's **consistency** — especially on moral/ethical questions — independent of correctness. It checks whether the model gives the **same answer to semantically equivalent prompts**.

Semantic inconsistency (Q8): defined as **providing semantically disparate answers to semantically equivalent prompts** — i.e. you rephrase the same question and the model's meaning-level answer changes. (NOT about factual errors, grammar, or single-prompt reasoning.)

Semantic Graph Entropy (Q9): a **higher** semantic graph entropy indicates the model is exhibiting **higher semantic INCONSISTENCY** — its answers across paraphrases are more spread out / disordered.

Temperature & consistency (Q10): if a developer lowers the generation temperature to reduce contradictory moral answers, the SaGE expectation is the **consistency score will remain largely unaffected** — inconsistency is a deeper property than sampling randomness, so turning down temperature doesn't fix it.

Accuracy vs consistency (Q11): an LLM's accuracy on standard commonsense-reasoning benchmarks **does NOT correlate with its moral consistency** — a model can score well on benchmarks yet be morally inconsistent. They measure different things.

SaGE quick-key: Q8 = semantically disparate answers to equivalent prompts; Q9 = higher entropy → higher inconsistency; Q10 = lowering temperature leaves consistency largely unaffected; Q11 = accuracy does NOT correlate with consistency.

6 One-line recap for the exam

- **Right to be forgotten** = delete-on-request (GDPR Art.17)
- **OECD 'robustness, security & safety'** = safe under normal use/misuse/adverse conditions
- **OECD accountability** = traceability of data & decisions
- **EU AI Act** = horizontal, 4 risk tiers
- **Banned (unacceptable)** = social scoring + cognitive manipulation
- **Deepfakes** must be labelled AI-generated
- **Deployer** = uses the AI under its own authority
- **Silent mitigation** = handled in background, no user action
- **SaGE** = consistency across paraphrases; high entropy = inconsistent
- **Accuracy ≠ moral consistency**; temperature doesn't fix it

Heads-up: the screenshots captured Q1–Q11 only; **Week 11 Q12 wasn't visible**, so it's not in these notes. If you can send that question I'll add it.